

Chordal Sparse Matrices and Semidefinite Programming

Theo Diamandis
tdiamand@mit.edu

Semidefinite Programs (SDPs) have PSD constraints

Primal Problem: variable $Z \in \mathbf{S}^n$

$$\begin{array}{ll}\text{maximize} & \mathbf{Tr}(GZ) \\ \text{s.t.} & \mathbf{Tr}(F_i Z) + c_i = 0, \quad i = 1, \dots, m \\ & Z \succeq 0,\end{array}$$

Dual problem: variable $x \in \mathbf{R}^m$

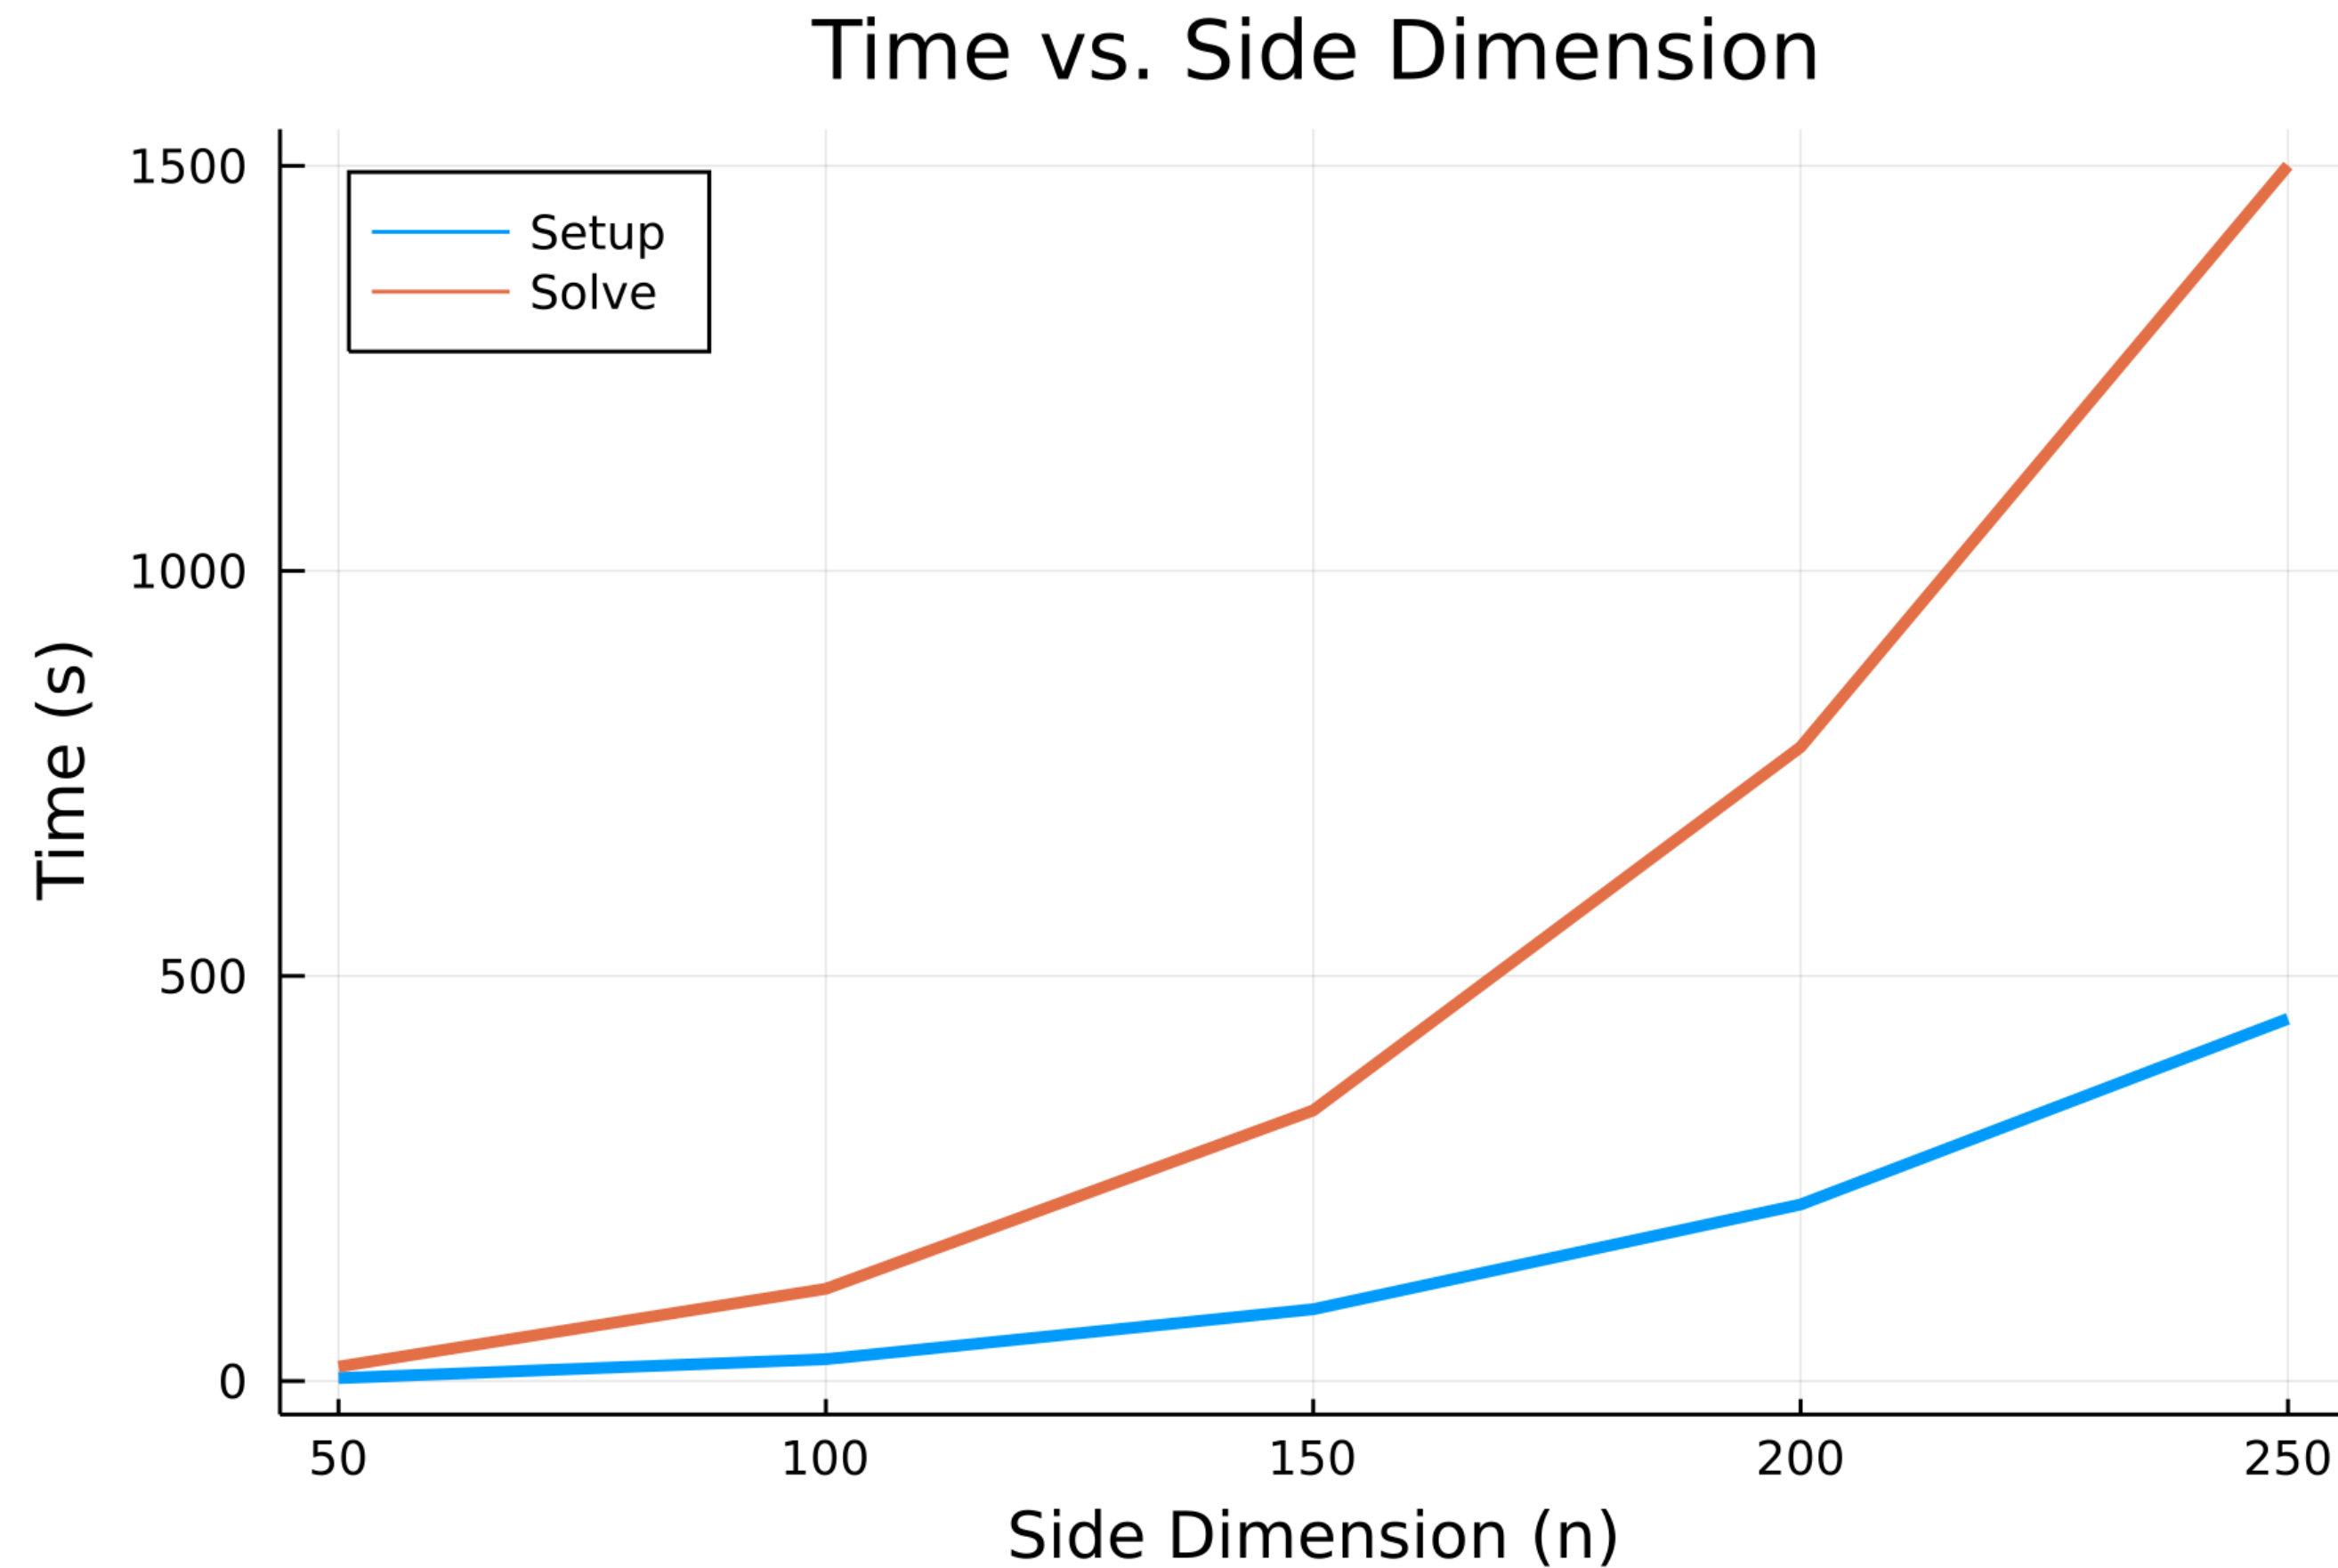
$$\begin{array}{ll}\text{minimize} & c^T x \\ \text{s.t.} & -\sum_i x_i F_i - G \succeq 0\end{array}$$

SDPs are a powerful modeling tool

- Signal Proc & Control
 - Phase Retrieval (e.g. imaging)
 - Polynomial controller design
- Stats/ML
 - Outlier Detection (e.g. ellipsoidal peeling)
 - Experiment Design
 - Factor Analysis
 - Low-rank matrix completion & decomp.
 - Neural network robustness verification
- Circuit Design
- Portfolio Optimization
- Combinatorial Optimization
- Mechanical structure optimization
- Operations Research
 - Facility Location
 - Scheduling
- Robust optimization
- Geometric data processing & CV

SDPs are plagued by scaling issues

Per iteration cost is (approx) cubic in the side dimension



But we can exploit sparsity

Variable $Z \in \mathbf{S}^n$

$$\begin{array}{ll} \text{maximize} & \text{Tr}(GZ) \\ \text{s.t.} & \text{Tr}(F_i Z) + c_i = 0, \quad i = 1, \dots, m \\ & Z \succeq 0, \end{array}$$


Only use, for example, 2% of entries of Z

} **Sparse**

Goal: get rid of the “dense” $n \times n$ PSD constraint \rightarrow reduce # variables by 98%

Need: a check that a sparse matrix can be “completed” to be a PSD matrix

Chordal Decomposition: check $Z \succeq 0$ via sub matrices




$O(5^3) = C \cdot 125$


$O(2 \cdot 3^3) = C \cdot 54$

Detour: Chordal Graph Theory


Sparsity patterns = undirected graphs



Sparsity patterns = undirected graphs



chordal



not chordal

Definition: A graph is *chordal* if all cycles of length > 3 have a chord

- Any graph can be made chordal by adding edges
- Many “hard” graph problems are “easy” on chordal graphs

Definition: A *clique* is a set of vertices that induces a maximal complete subgraph

Examples (Code Demo)

Example sparsity patterns:

- Banded
 - Arrow
 - Random
-
- Think about which are Chordal

Check if Z is PSD by looking at sub matrices

[Grone et al. 1984]

Theorem: A sparse matrix Z with a chordal sparsity pattern can be PSD completed if and only if

$$Z_{C_p} \succeq 0, \quad p = 1, \dots, P$$

where Z_{C_p} is the sub matrix of Z indexed by the clique C_p .

Applied to SDPs, reduces PSD constraint size

Matrices F_i , G have a chordal aggregate sparsity pattern

$$\begin{array}{ll} \text{maximize} & \text{Tr}(GZ) \\ \text{s.t.} & \text{Tr}(F_i Z) + c_i = 0, \quad i = 1, \dots, m \\ & \cancel{Z \succeq 0}, \quad \boxed{Z_{C_p} \succeq 0}, \end{array}$$

Iteration cost is now $O\left(\sum_{p=1}^P |C_p|^3\right)$ vs $O(n^3)$

See COSMO.jl & Chordal.jl

Use similar approach for the dual problem

[Agler et al. 1988]

Theorem: A sparse matrix S with a **chordal sparsity pattern** is PSD if and only if there exist matrices $S_p \succeq 0$ such that

$$S = \sum_{p=1}^P T_p S_p T_p^T$$

where T_p is a selector matrix:

$$[T_p]_{ij} = \begin{cases} 1 & C_p(i) = j, \\ 0 & \text{otherwise} \end{cases}$$

Use a similar approach for the dual problem

Matrices F_i , G have a chordal aggregate sparsity pattern

$$\begin{array}{ll} \text{minimize} & c^T x \\ \text{s.t.} & -\sum_i x_i F_i - G \end{array}$$

$$= \sum_{p=1}^P T_p S_p T_p^T$$

$$S_p \succeq 0$$

Algorithmic Details

These constructions lead to two new cones:

Sparse PSD-completable cone (primal problem):

$$\Pi_E(S_+^n) = \{X \mid \exists Y \in S_+^n \text{ s.t. } X = \Pi_E(Y)\}$$

Sparse PSD cone (dual problem):

$$S_+^n \cap S_E^n = \{X \mid X \succeq 0 \text{ and } (i,j) \in E \implies X_{ij} = 0\}$$

- We want to compute projections and gradients/Hessians of barriers

Reformulating \rightarrow faster FO & IP algorithms

- For first order algorithms (e.g., ADMM), we project onto the PSD cone
 - Chordal decomposition \rightarrow break up large PSD projection into P small ones

$$X^{k+1} = \operatorname{argmin}_X \left(f(X) + (\rho/2) \|X - Z^k + U^k\|_F^2 \right)$$




$$Z^{k+1} = \Pi_{S_+^n} \left(X^{k+1} + U^k \right)$$

$$U^{k+1} = U^k + X^{k+1} - Z^{k+1}.$$

Reformulating \rightarrow faster FO & IP algorithms

- For first order algorithms (e.g., ADMM), we project onto the PSD cone
 - Chordal decomposition \rightarrow break up large PSD projection into P small ones
- For interior point algorithms, must compute gradient and/or Hessian of barrier
 - Chordal decomposition \rightarrow efficiently compute these [ADV13, VA15]

Chordal Decomposition Steps



1

Input Sparsity Pattern

2

Chordal Extension

Reordering matters!

3

Find (& merge) cliques

[Michael Garstka et al. 2020]

4




Form SDP

5

Solve SDP

6

[primal] Reconstruct Z



We can reconstruct the full Z if needed

Our optimization variable Z is only partially specified

Two main approaches:

1. Maximum determinant completion [VA15]

- Idea: Solve $ZL = L^{-T}D^{-1}$, where factorization is known from decomp

2. Minimum rank completion [Sun15]

- Idea: factor each clique block & apply necessary orthogonal rotations

$$r = \max_p \mathbf{rank}(Z_{C_p})$$

Numerical Example

Example: Maximum Variance Unfolding

Preserve local distances + maximize (global) variance

Data: $x_i \in \mathbb{R}^{d_x}$ with pairwise distances $D_{ij} = \|x_i - x_j\|^2$

Goal: find embedding $y_i \in \mathbb{R}^{d_y}$ that preserves *local* distances

Method: solve an SDP with variable $Y = \tilde{Y}^T \tilde{Y}$, where

$$\tilde{Y} = \begin{bmatrix} | & & | \\ y_1 & \dots & y_n \\ | & & | \end{bmatrix} \in \mathbb{R}^{d_y \times n}$$

maximize

$\text{Tr } Y$

← Variance

s.t.

$$Y_{ii} + Y_{jj} - 2Y_{ij} = D_{ij}$$


$$\forall (i, j) \in \mathcal{E}_k$$

$$\sum y_i = 0$$

$$\mathbf{1}^T Y \mathbf{1} = 0$$

$$Y \succeq 0.$$

← Preserve kNN distances



MVU Sparsity Pattern

$N = 250$, $k = 6$


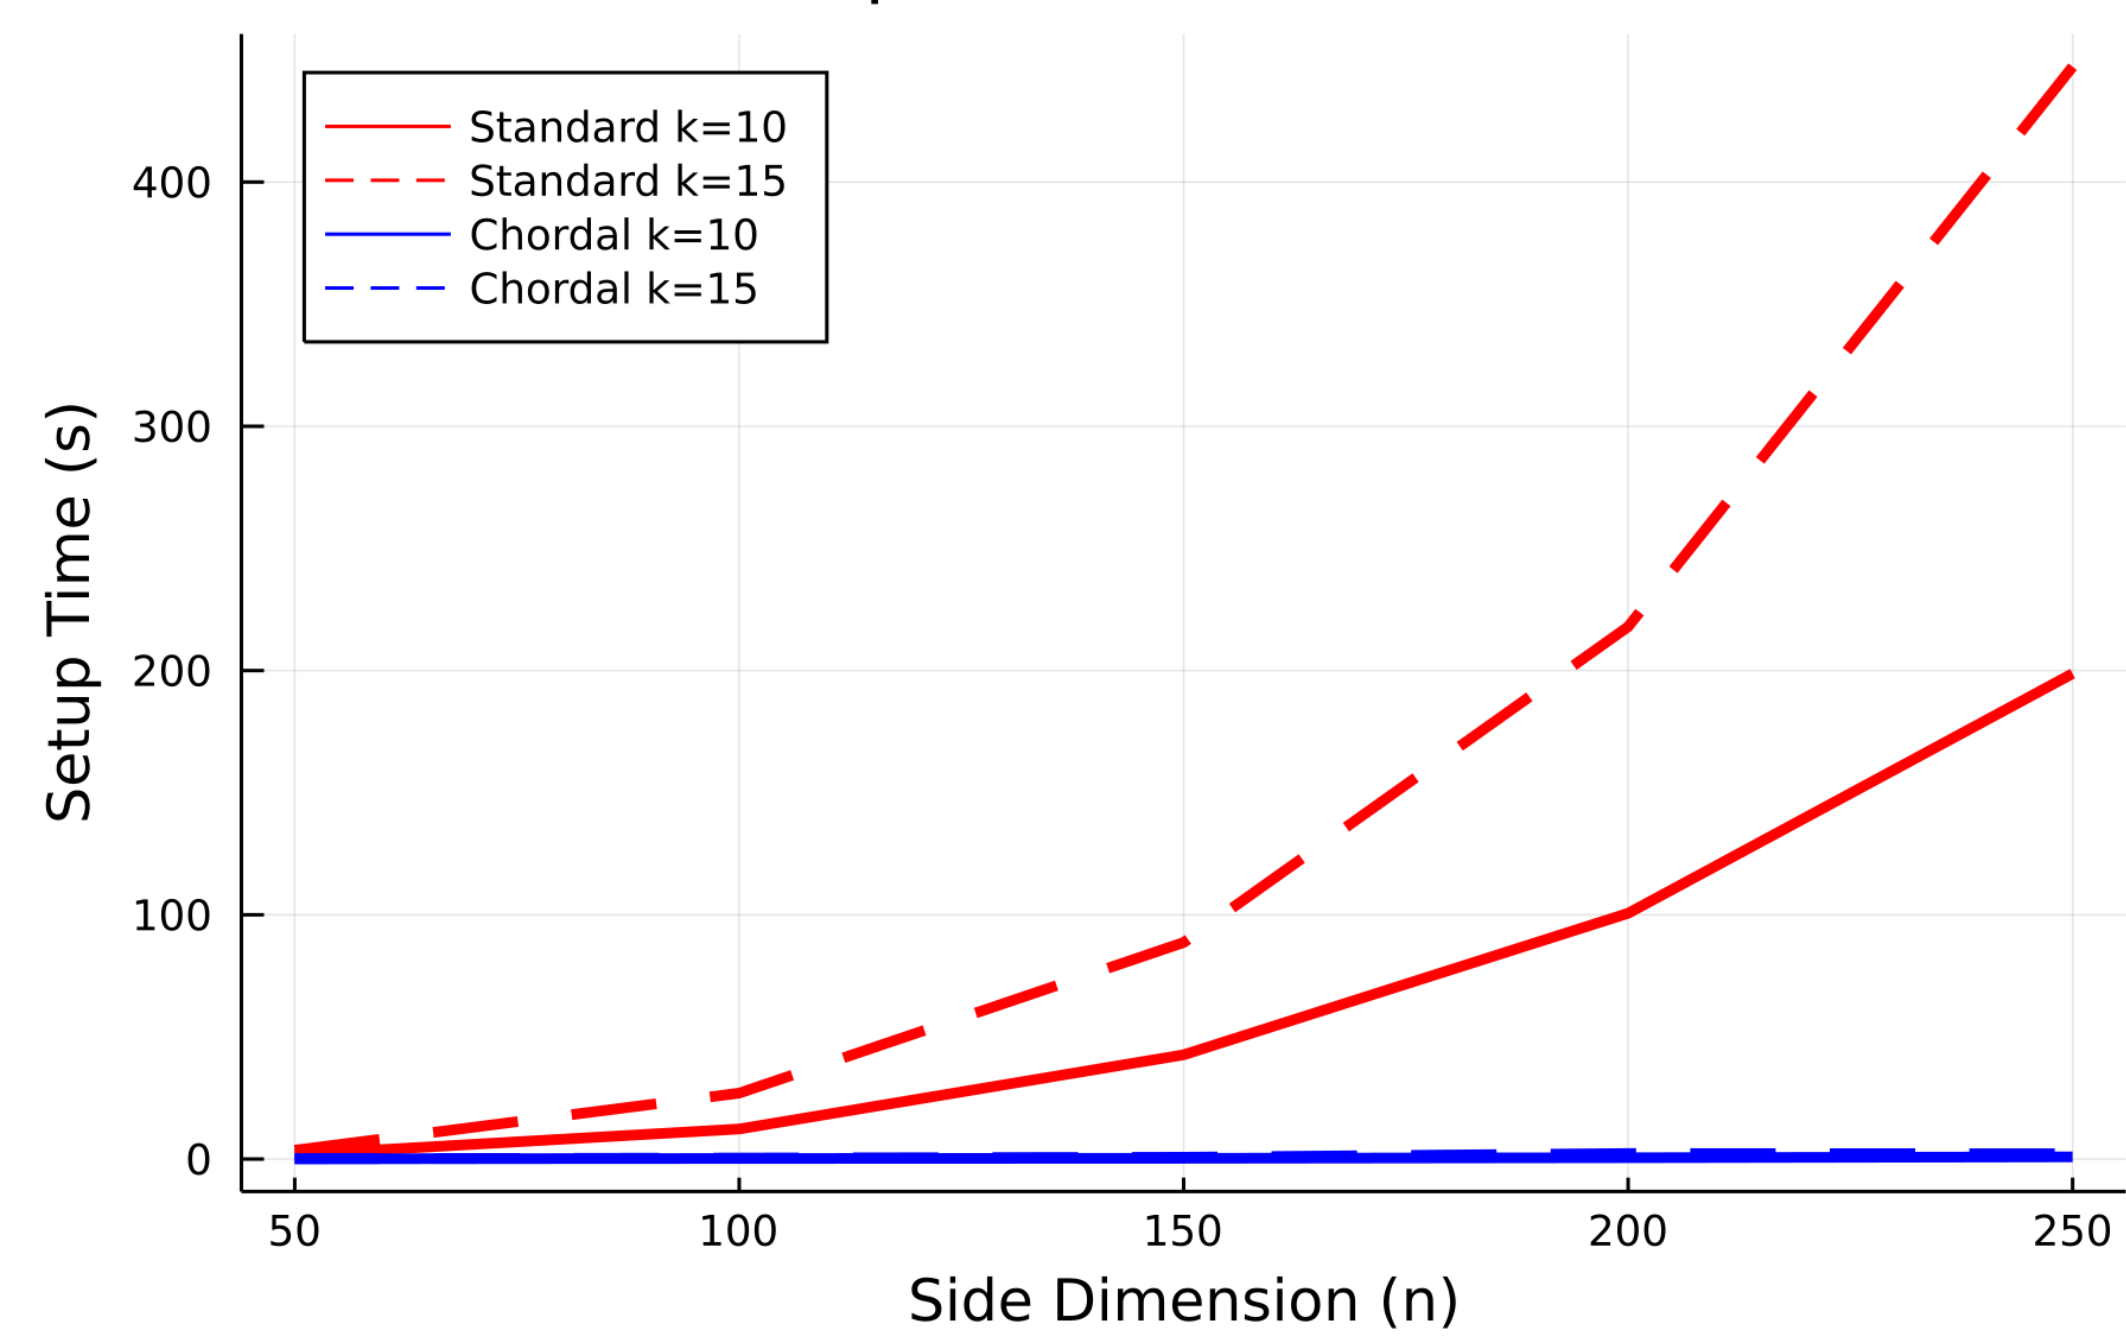



Figure 3: Aggregate sparsity pattern for MVU with $n = 250$ and $k = 6$, which has 6487 nonzeros (left) and its chordal extension, which has 7382 nonzeros (right).

Chordal decomposition *dramatically* speeds up SDPs

SDP Setup Time vs. Side Dimension


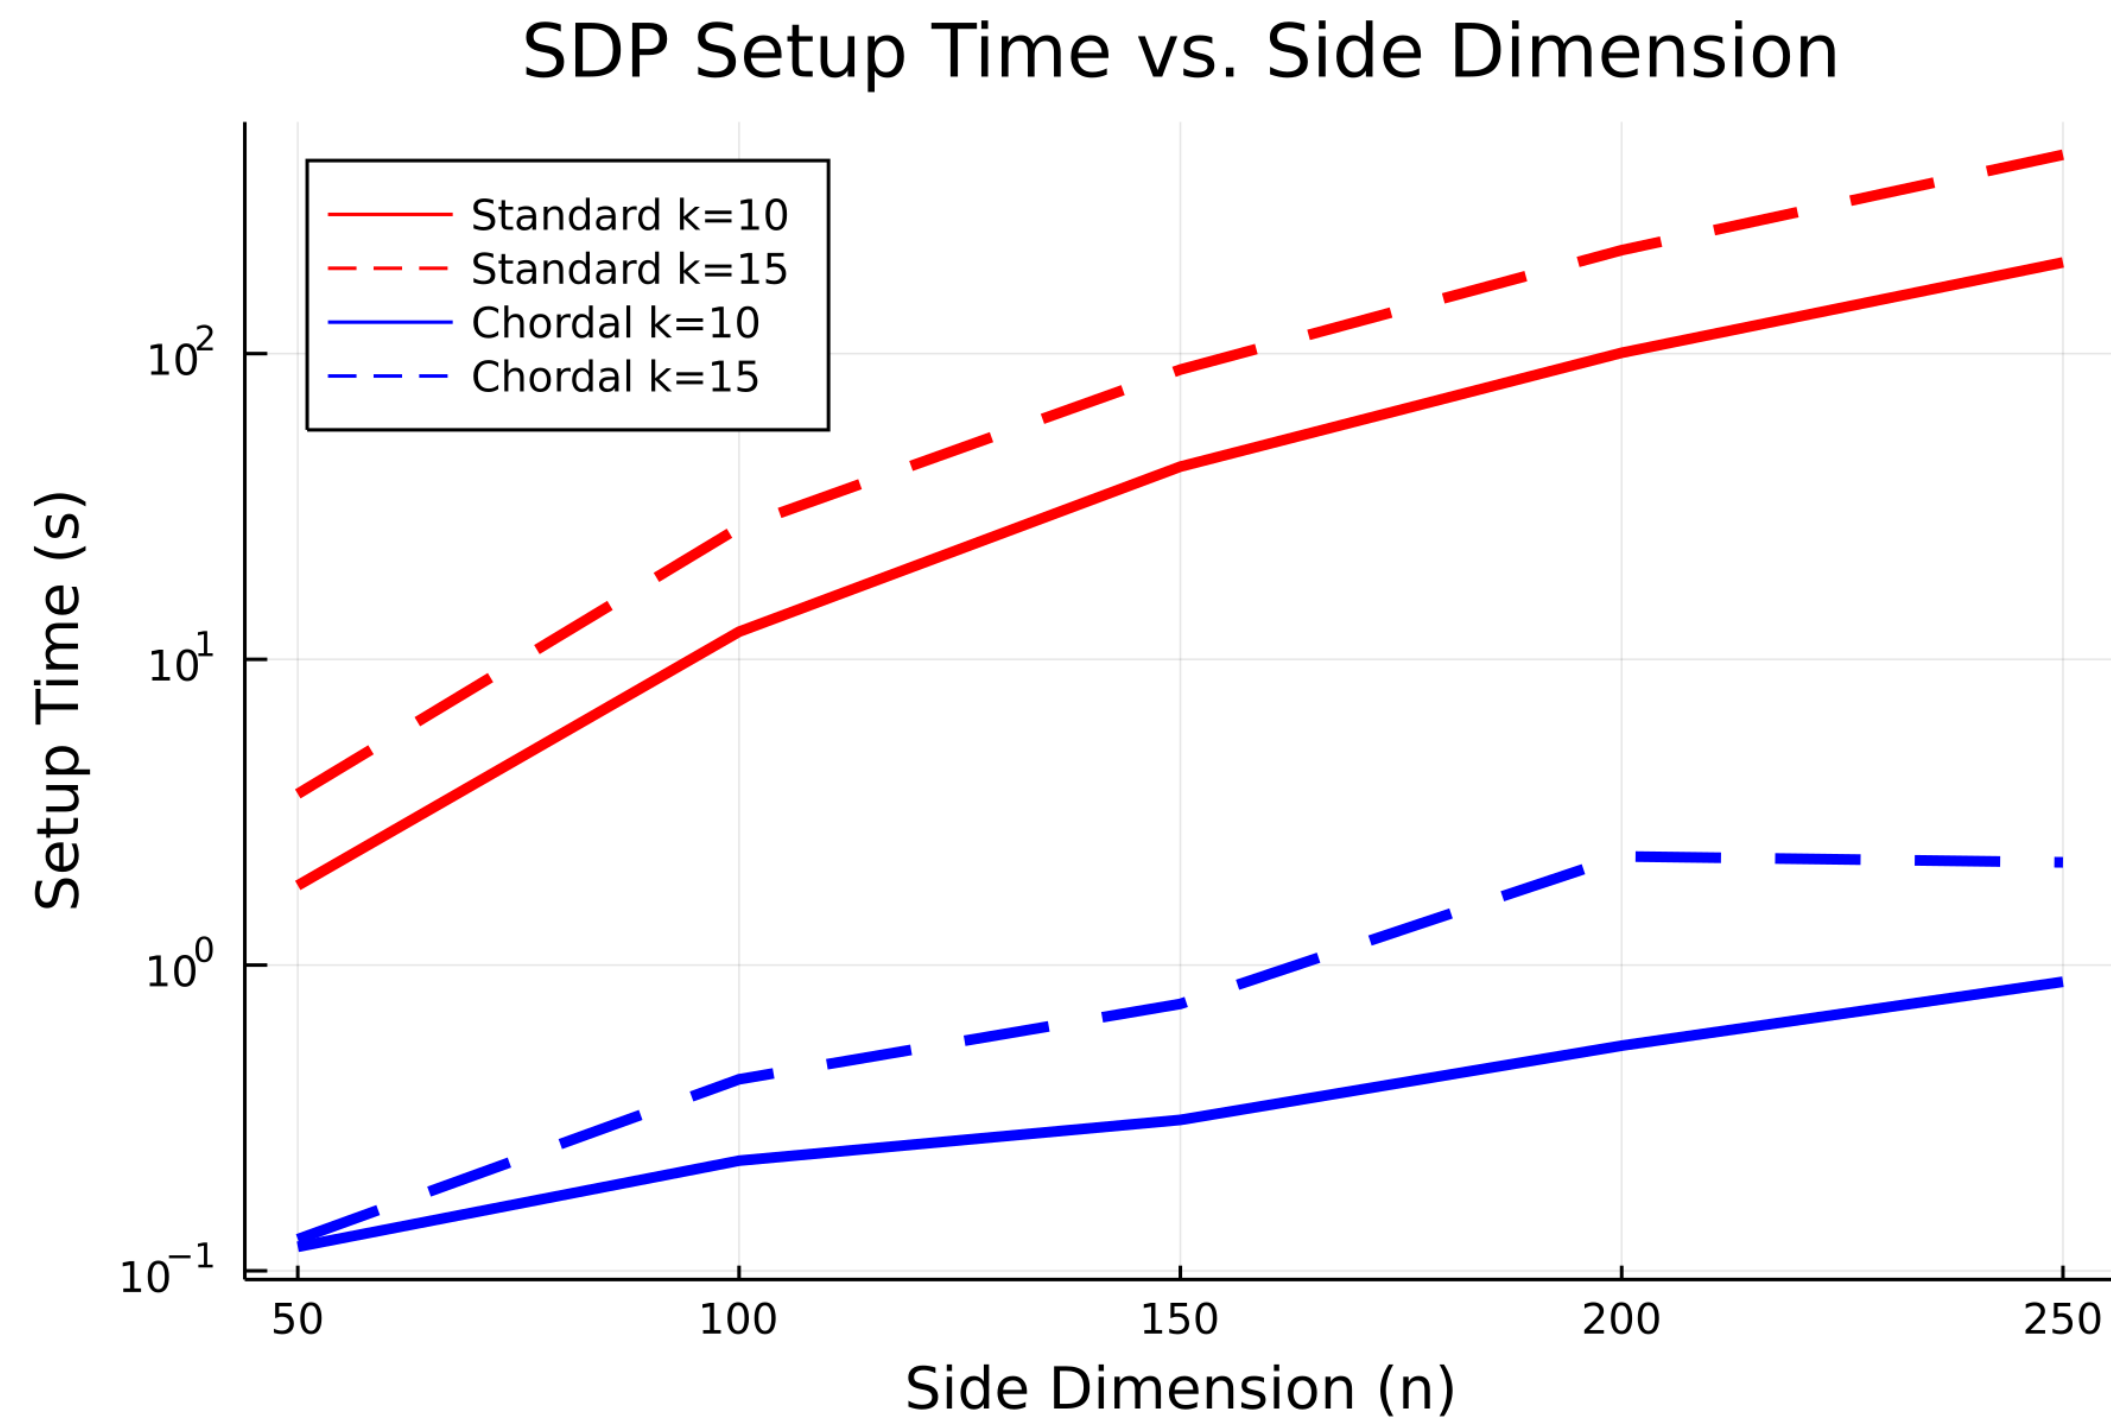


SDP Solve Time vs. Side Dimension



Up to 175x faster on maximum variance unfolding!

Chordal decomposition *dramatically* speeds up SDPs



Up to 175x faster on maximum variance unfolding!

Wrap up

Chordal.jl has several functions for sparse matrices

Survey: [Vandenberghe and Andersen 2015]

- Chordal decomposition of SDPs (with JuMP.jl) [Fukuda et al. 2001]
- Elimination trees, clique trees
- Clique graphs & merging [Garstka et al. 2020]
- Maximum determinant & minimum rank PSD completion [Sun 2015]
- Chordality tests [Tarjan and Yannakakis 1984]
- Euclidean distance matrix completion
- Several examples

Project Ideas

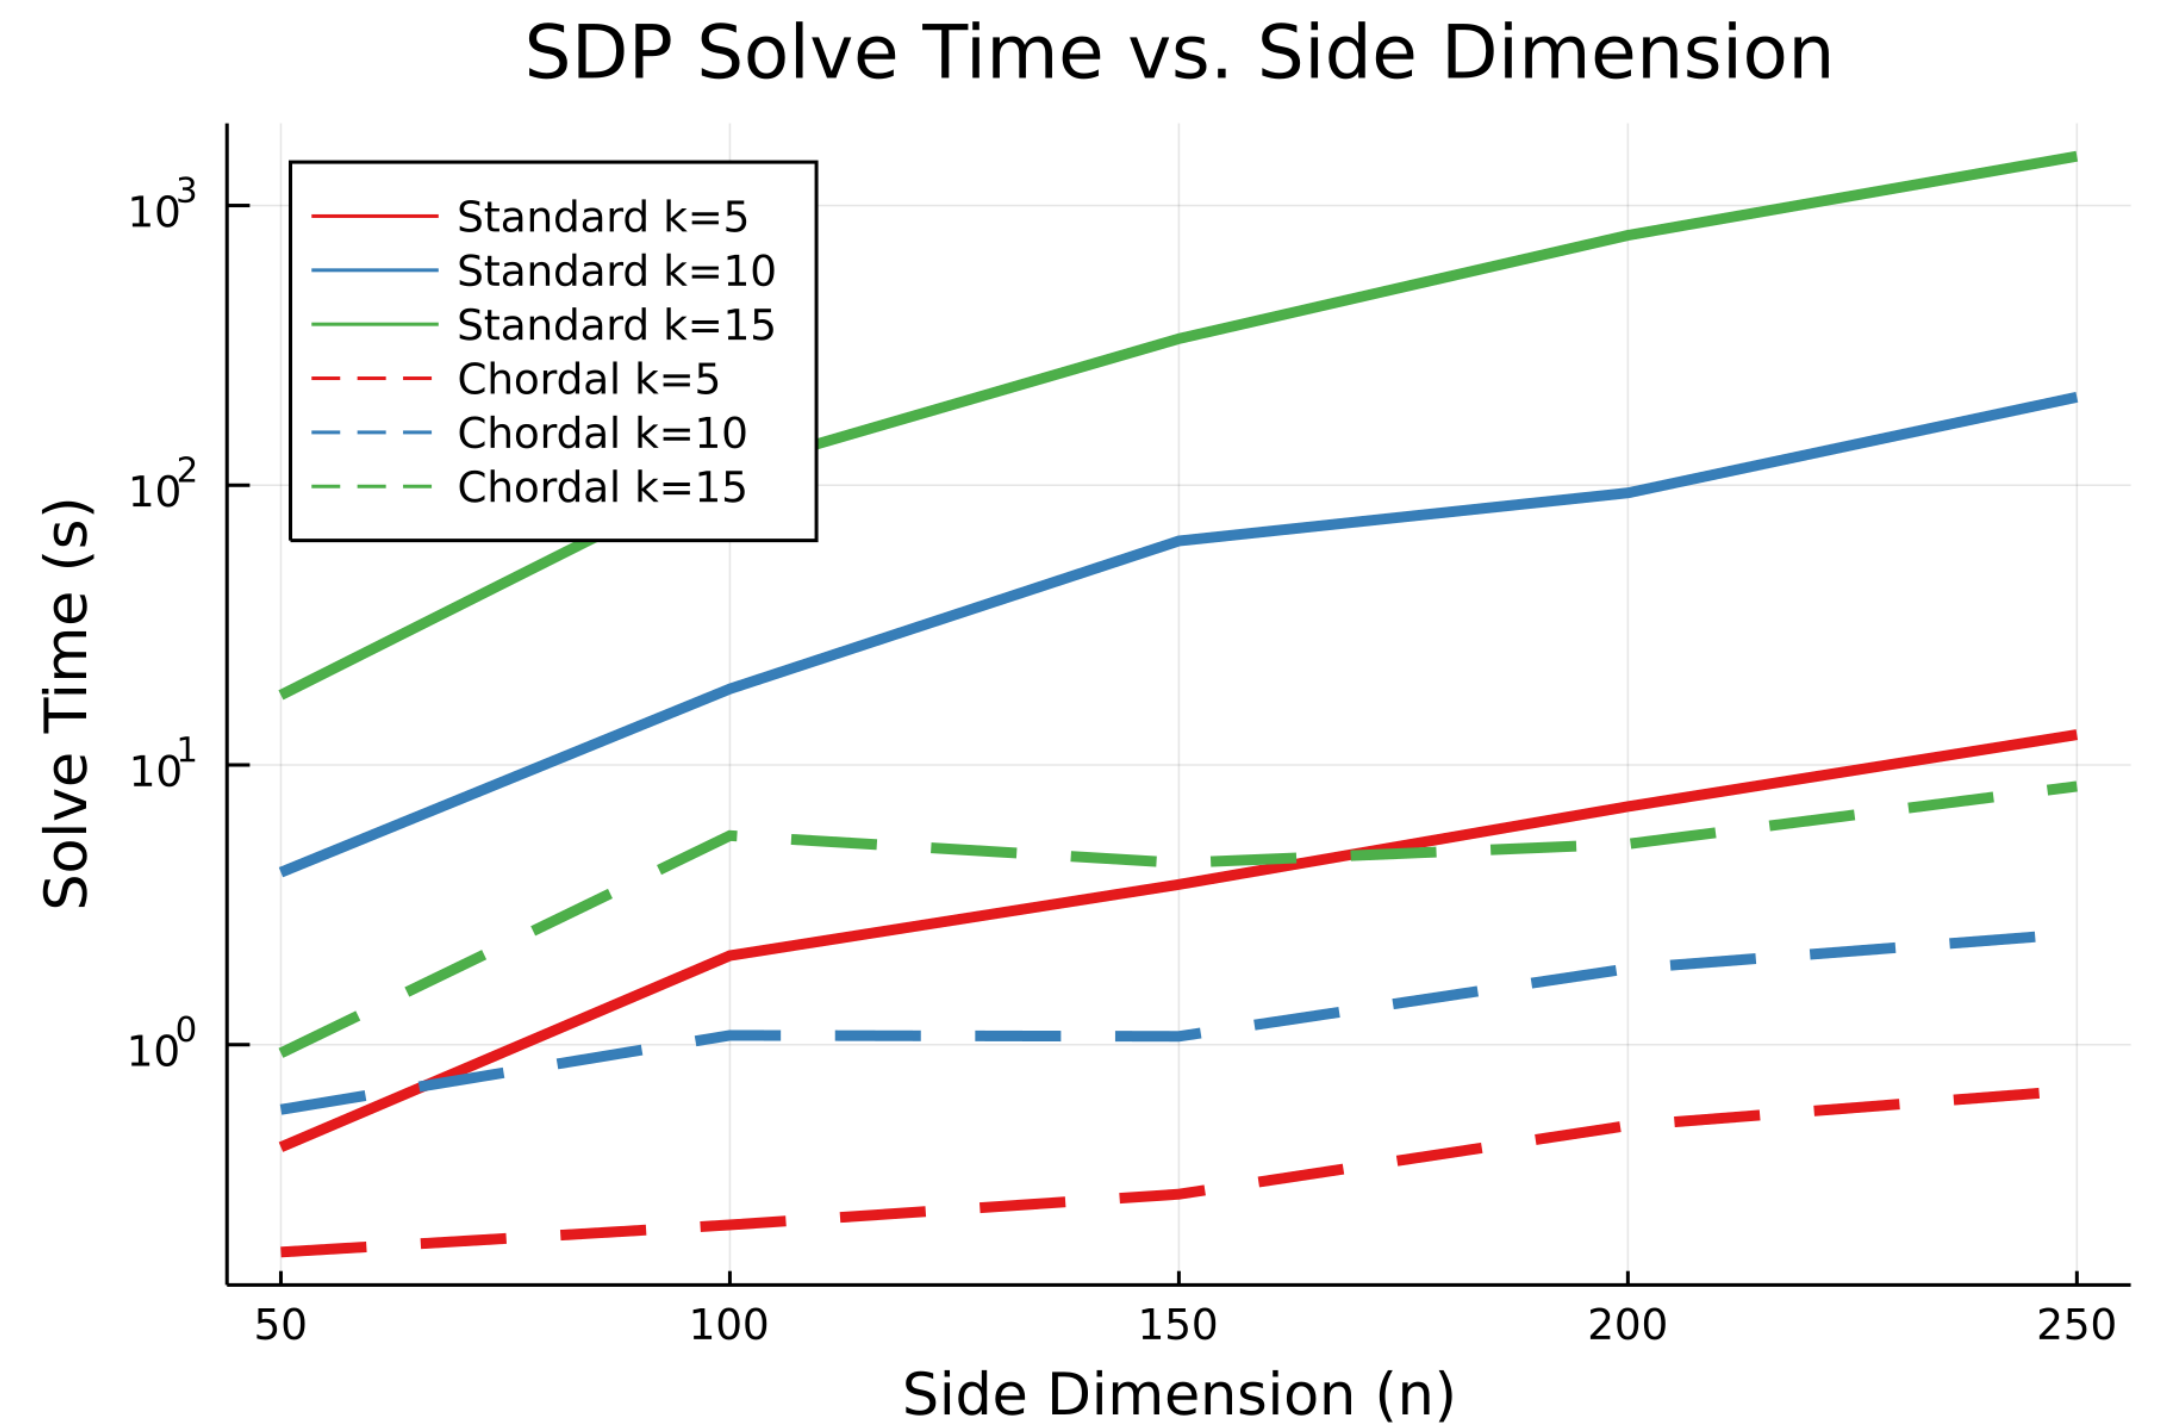
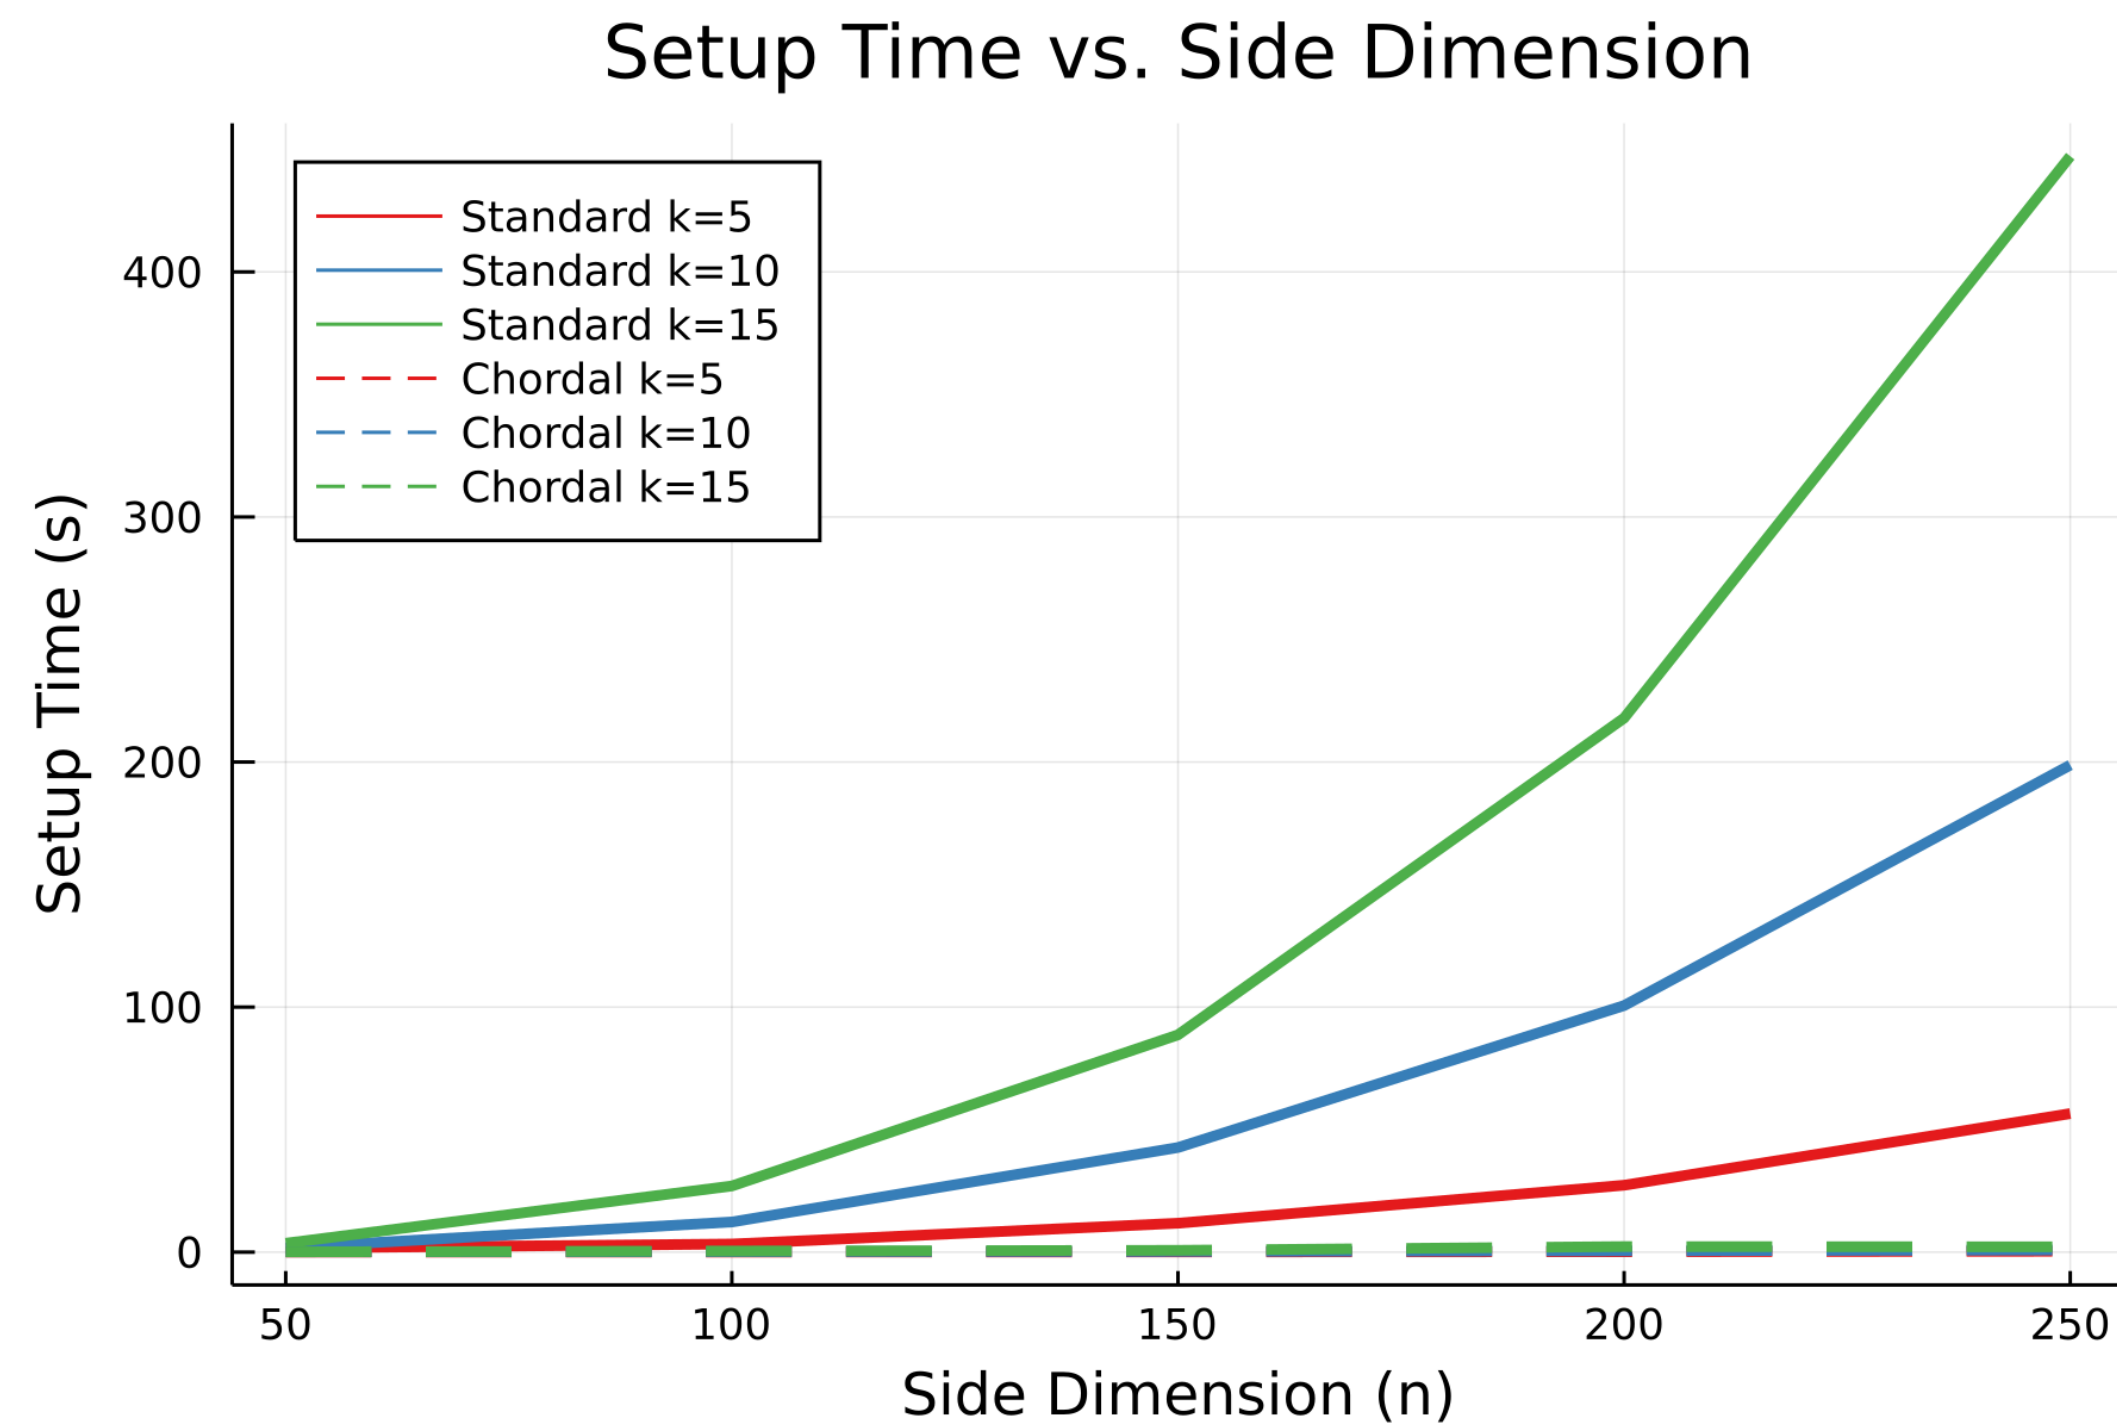
- Use Maximum Variance Unfolding to embed and then classify MNIST
 - Q: Does error propagate up the clique tree? Can we see this in accuracy?
- Investigate scaling of decomposed vs. standard sparse SDPs
 - Big issue in SDP solvers (especially first order ones)
- Implement a pure Julia sparse super nodal Cholesky
- Use these techniques for discrete optimization problems (see Ch 7 in [VA15])
- Integrate with Jump.jl (maybe not for this class, but great open source contribution!)

References




- Martin Andersen et al. “Logarithmic barriers for sparse matrix cones.” In: Optimization Methods and Software 28.3 (2013), pp. 396–423
- Mituhiro Fukuda et al. “Exploiting sparsity in semidefinite programming via matrix completion I: General framework.” In: SIAM Journal on Optimization 11.3 (2001), pp. 647–674.
- Michael Garstka, Mark Cannon, and Paul Goulart. “A clique graph based merging strategy for decomposable SDPs.” In: IFAC-PapersOnLine 53.2 (2020), pp. 7355–7361.
- Yifan Sun. “Decomposition methods for semidefinite optimization.” PhD thesis. UCLA, 2015.
- Robert E Tarjan and Mihalis Yannakakis. “Simple linear-time algorithms to test chordality of graphs, test acyclicity of hypergraphs, and selectively reduce acyclic hypergraphs.” In: SIAM Journal on computing 13.3 (1984), pp. 566–579.
- Lieven Vandenberghe and Martin S Andersen. “Chordal graphs and semidefinite optimization.” In: Foundations and Trends in Optimization 1.4 (2015), pp. 241–433.

Appendix

Chordal decomposition *dramatically* speeds up SDPs (Primal form SDP)



Up to 200x faster on maximum variance unfolding!



Definition 4.1 (Running intersection property).

For each pair of cliques $\mathcal{C}_i, \mathcal{C}_j \in \mathcal{B}$, the intersection $\mathcal{C}_i \cap \mathcal{C}_j$ is contained in all the cliques on the path in the clique tree connecting \mathcal{C}_i and \mathcal{C}_j .

This property is also referred to as the *clique-intersection property* in [NFF⁺03] and the *induced subtree property* in [VA⁺15]. For a given chordal graph, a clique tree can be computed using the algorithm described in [PS90]. The clique tree for an example sparsity pattern is shown in Figure 1(c).

In a *post-ordered* clique tree the descendants of a node are given consecutive numbers, and a suitable post-ordering can be found via depth-first search. For a clique \mathcal{C}_ℓ we refer to the first clique encountered on the path to the root as its *parent clique* \mathcal{C}_{par} . Conversely \mathcal{C}_ℓ is called the *child* of \mathcal{C}_{par} . If two cliques have the same parent clique we refer to them as *siblings*.